

# How to Secure OpenShift Environments and What Happens If You Don't


JAN HARRIE

## \$ whoami – Jan Harrie




- Security Consultant @ERNW GmbH
- Former Security Analyst/Pentester/WebApp-Monkey/Social-Engineer
- M.Sc. IT-Security TU Darmstadt
- Research // Interests:
  - K8s on-prem solutions
  - Cluster extensions
  - Gardening



 @NodyTweet

# Agenda

1. OpenShift & Kubernetes – Introduction & Differences
2. Cluster Threats
3. (In-)Security of Clusters
4. Conclusion & Future Work





# OpenShift & Kubernetes – Introduction & Differences

# Introduction OpenShift

- (On-Premise) Container Execution Platform from RedHat
- First Release 05/2011
- Current Stable Release: 4.2 (11/2019)
- Host Operation System is RedHat Enterprise Linux and Container Linux from CoreOS
- Since Version 3 with K8s under the hood
- Since Version 4 Based on CRI-O, previously Docker
- OKD Community Version, e.g., CentOS
  - Current Stable Release (10/2018): v3.11
  - Builds on K8s 1.11

# OpenShift vs. K8s – Differences



## Kubernetes

- Role Based Access Control
- Namespaces
- Resource Limits
- Security Context
- Network Policies
- Pod Security Policies



## OpenShift

- Image Streams
- Application Catalogue
- User Management
- Templates
- Revision History
- Security Context Constraints



# Cluster Threats

## Cluster Threats


**What kind of threat model exist  
for a cluster?**



# Cluster Threats

## External Attacker


- (Only) Access to Offered Services
- No API Access
- No Cluster-Insights Knowledge
- Maybe public knowledge from DockerHub and Quay or GitHub




# Cluster Threats

## Internal Attacker


- API Access
- Control over Images and Deployments
- Access to Code Repositories
- Internal Cluster Knowledge



# Internal are External Attackers one Step ahead



# Internal are External Attackers one Step ahead




## Internal Attacker

- API Access
- Control over Images and Deployments
- Access to Code Repositories
- “Cluster Internal Knowledge”



# (In-)Security of Clusters





**Container Escapes**

**Kernel Level Exploits**

**Seccomp**

**SELinux**

**Pod Security Policies**

**Image Safety**

**RBAC**

**Network Isolation**


**Low Hanging  
Fruit**

**Logging**

**Least Privilege**

**OS Hardening**





Container Escapes

Kernel Level Exploits

Seccomp

SELinux

Pod Security Policies

Image Safety

RBAC

Network Isolation

Low Hanging  
Fruit

Logging

Least Privilege

OS Hardening



# (In-)Security of the Cluster

User Management

Network Security

A dive into Security Context Constraints (SCC's)



# (In-)Security of the Cluster

**User Management**

Network Security

A dive into Security Context Constraints (SCC's)

# User Management in OpenShift

## OpenShift offers integration into multiple Identity Provider (IdP)

- E.g., HTTPasswd, Keystone, LDAP authentication, Basic authentication (remote), Request header, GitHub, GitLab, Google, OpenID connect ; one IdP configurable
- Implicit: `mappingMethod: claim`, Explicit: `mappingMethod: lookup`

## Identities are Mapped to User in the Cluster

- Identities are based on the IdP, while a User is an Object in the Cluster

## Users can be organized in Groups

- LDAP sync and manual assignment possible

## “True User Removal” only possible in the IdP

- Manual deleted Users and Identities are re-created on next login.

# Role Based Access Control

A lot of default cluster-roles are shipped with OpenShift

- Introduction of new roles is recommended rather than adjustment
- Modification may lead to broken functionality

## Authenticated User:

- Implicit association with virtual group `system:authenticated` // `system:authenticated:oauth`
- What does this mean?

Demo 1: <https://asciinema.org/a/281016>

# What can probably go wrong?

If IdP Is wrong configured:

→ Users can deploy workload in the cluster

and

→ Inspection of `resolv.conf` of the Pods:

```
$ cat /etc/resolv.conf
```

```
nameserver 172.30.0.2
```

```
search user1-p0.svc.cluster.local svc.cluster.local cluster.local
```

```
options ndots:5
```

Demo 2: <https://asciinema.org/a/281704>

# Mitigation Strategy

- Patch the Cluster Role:

```
$ oc adm policy remove-cluster-role-from-group self-provisioner system:authenticated  
clusterrolebinding.rbac.authorization.k8s.io/self-provisioners patched
```

```
$ oc login -u user1
```

```
$ oc new-project user1-p1
```

```
Error from server (Forbidden): You may not request a new project via this API.
```

- Define DNS policy per Pod [12]

# (In-)Security of Clusters

User Management

## **Network Security**

A dive into Security Context Constraints (SCC's)

# Network Security

## Software Defined Networking build on Open vSwitch

Three plugins available:

- Open vSwitch Subnet
  - Open vSwitch Multitenant
  - Open vSwitch Networkpolicy
- Master-Nodes do not participate in the Cluster Network
  - Each Node gets its own Class-C network for the Pods assigned
  - Overlay communication via VXLAN
  - Integration of other Hosts into the cluster network by:
    - Host as an OpenShift node
    - Creating a VXLAN tunnel

# Network Security – Open vSwitch Subnet

Configuration of `Open vSwitch Subnet` is not recommended

→ Cross project communication is possible

Demo 3: <https://asciinema.org/a/280323>



# Network Security – Open vSwitch Multitenant

Setup Plugin **Open vSwitch Multitenant** to “prevent” cross-project communication

- Each Project get its own Virtual Network ID (VNID)
- Communication between different projects prohibit.
- Projects can be joined together

**BUT !**

- Separation on Namespace-Level
- Projects with ***VNID 0*** are more privileged
- The project ***default*** has ***VNID 0***

Side reference: TR19 – VXLAN Security or Injection [8]

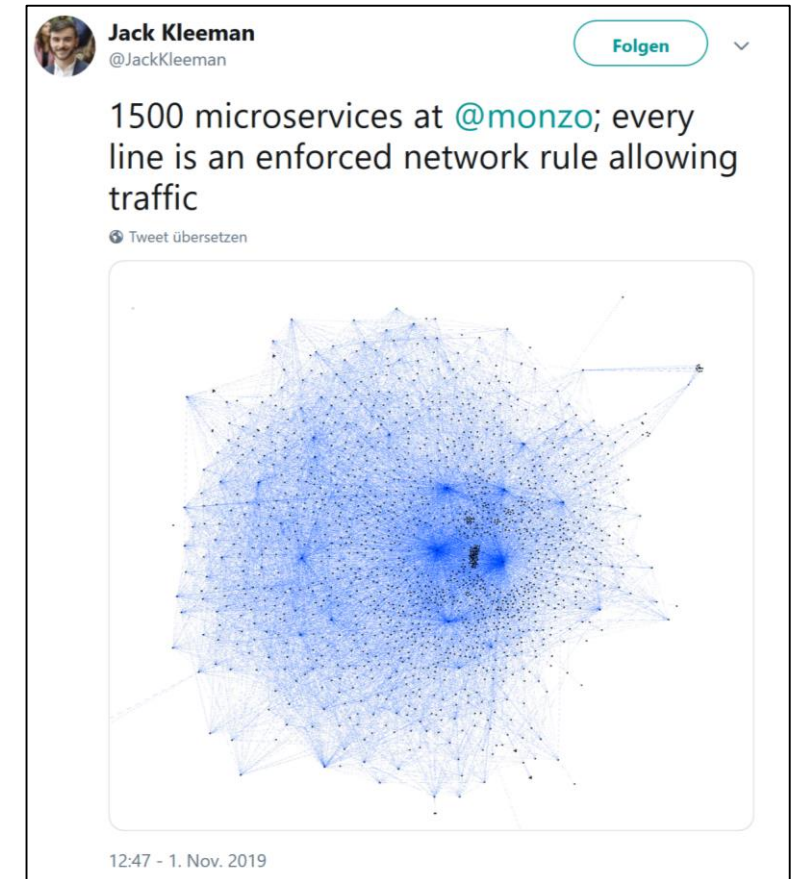
# Network Security – Open vSwitch Networkpolicy

Alternatively: stick to ovs-networkpolicy which allows you to deploy NetworkPolicies, and block all ingress traffic [9] and add explicit whitelisting.

```
kind: NetworkPolicy
metadata:
  name: default-deny
spec:
  podSelector: {}
  policyTypes:
  - Ingress
```

Further more, the plugin allows White- an Black-Listing on Layer3 [10] with CIDR notation or DNS

Configuration of Egress IP's and Egress Proxies is possible [11]



<https://twitter.com/JackKleeman/status/1190354757308862468>

# (In-)Security of Clusters

User Management

Network Security

**A dive into Security Context Constraints (SCC's)**

# Security Context Constraints

- Introduced by release 3.0 (05/2015)
- Secure Context Constraints (SCC's) is for Pods what RBAC is for the SAs
- Restrict execution of Pods
- Created by Cluster Administrator and assigned to Service Account
- Default SCC is 'restricted'

# Security Context Constraints

## Predefines Profiles

```
$ oc get scc
```

| NAME             | PRIV  | CAPS | SELINUX   | RUNASUSER        | [...] |
|------------------|-------|------|-----------|------------------|-------|
| anyuid           | false | []   | MustRunAs | RunAsAny         | [...] |
| hostaccess       | false | []   | MustRunAs | MustRunAsRange   | [...] |
| hostmount-anyuid | false | []   | MustRunAs | RunAsAny         | [...] |
| hostnetwork      | false | []   | MustRunAs | MustRunAsRange   | [...] |
| nonroot          | false | []   | MustRunAs | MustRunAsNonRoot | [...] |
| privileged       | true  | [*]  | RunAsAny  | RunAsAny         | [...] |
| restricted       | false | []   | MustRunAs | MustRunAsRange   | [...] |

# Security Context Constraints

Predefines Profiles – that allow **privileged**

```
$ oc get scc
```

| NAME              | PRIV        | CAPS       | SELINUX         | RUNASUSER        | [...]        |
|-------------------|-------------|------------|-----------------|------------------|--------------|
| anyuid            | false       | []         | MustRunAs       | RunAsAny         | [...]        |
| hostaccess        | false       | []         | MustRunAs       | MustRunAsRange   | [...]        |
| hostmount-anyuid  | false       | []         | MustRunAs       | RunAsAny         | [...]        |
| hostnetwork       | false       | []         | MustRunAs       | MustRunAsRange   | [...]        |
| nonroot           | false       | []         | MustRunAs       | MustRunAsNonRoot | [...]        |
| <b>privileged</b> | <b>true</b> | <b>[*]</b> | <b>RunAsAny</b> | <b>RunAsAny</b>  | <b>[...]</b> |
| restricted        | false       | []         | MustRunAs       | MustRunAsRange   | [...]        |

Demo 4: <https://asciinema.org/a/280383>

# Security Context Constraints

Predefines Profiles – that allow **hostPath**, **hostIPC**, **hostPID**

```
$ oc get scc
```

| NAME                    | PRIV         | CAPS       | SELINUX          | RUNASUSER             | [...]        |
|-------------------------|--------------|------------|------------------|-----------------------|--------------|
| anyuid                  | false        | []         | MustRunAs        | RunAsAny              | [...]        |
| <b>hostaccess</b>       | <b>false</b> | <b>[]</b>  | <b>MustRunAs</b> | <b>MustRunAsRange</b> | <b>[...]</b> |
| <b>hostmount-anyuid</b> | <b>false</b> | <b>[]</b>  | <b>MustRunAs</b> | <b>RunAsAny</b>       | <b>[...]</b> |
| hostnetwork             | false        | []         | MustRunAs        | MustRunAsRange        | [...]        |
| nonroot                 | false        | []         | MustRunAs        | MustRunAsNonRoot      | [...]        |
| <b>privileged</b>       | <b>true</b>  | <b>[*]</b> | <b>RunAsAny</b>  | <b>RunAsAny</b>       | <b>[...]</b> |
| restricted              | false        | []         | MustRunAs        | MustRunAsRange        | [...]        |

Demo 5: <https://asciinema.org/a/280388>

# Security Context Constraints

Predefines Profiles – that allow **root in container**

```
$ oc get scc
```

| NAME                    | PRIV         | CAPS       | SELINUX          | RUNASUSER        | [...]        |
|-------------------------|--------------|------------|------------------|------------------|--------------|
| <b>anyuid</b>           | <b>false</b> | <b>[]</b>  | <b>MustRunAs</b> | <b>RunAsAny</b>  | <b>[...]</b> |
| hostaccess              | false        | []         | MustRunAs        | MustRunAsRange   | [...]        |
| <b>hostmount-anyuid</b> | <b>false</b> | <b>[]</b>  | <b>MustRunAs</b> | <b>RunAsAny</b>  | <b>[...]</b> |
| hostnetwork             | false        | []         | MustRunAs        | MustRunAsRange   | [...]        |
| nonroot                 | false        | []         | MustRunAs        | MustRunAsNonRoot | [...]        |
| <b>privileged</b>       | <b>true</b>  | <b>[*]</b> | <b>RunAsAny</b>  | <b>RunAsAny</b>  | <b>[...]</b> |
| restricted              | false        | []         | MustRunAs        | MustRunAsRange   | [...]        |



# Security Context Constraints

## Predefines Profiles – available

```
$ oc get scc
```

| NAME               | PRIV         | CAPS      | SELINUX          | RUNASUSER             | [...]        |
|--------------------|--------------|-----------|------------------|-----------------------|--------------|
| anyuid             | false        | []        | MustRunAs        | RunAsAny              | [...]        |
| hostaccess         | false        | []        | MustRunAs        | MustRunAsRange        | [...]        |
| hostmount-anyuid   | false        | []        | MustRunAs        | RunAsAny              | [...]        |
| <b>hostnetwork</b> | <b>false</b> | <b>[]</b> | <b>MustRunAs</b> | <b>MustRunAsRange</b> | <b>[...]</b> |
| nonroot            | false        | []        | MustRunAs        | MustRunAsNonRoot      | [...]        |
| privileged         | true         | [*]       | RunAsAny         | RunAsAny              | [...]        |
| restricted         | false        | []        | MustRunAs        | MustRunAsRange        | [...]        |

# Security Context Constraints

## Predefines Profiles – available

```
$ oc get scc
```

| NAME             | PRIV  | CAPS | SELINUX   | RUNASUSER        | [...] |
|------------------|-------|------|-----------|------------------|-------|
| anyuid           | false | []   | MustRunAs | RunAsAny         | [...] |
| hostaccess       | false | []   | MustRunAs | MustRunAsRange   | [...] |
| hostmount-anyuid | false | []   | MustRunAs | RunAsAny         | [...] |
| hostnetwork      | false | []   | MustRunAs | MustRunAsRange   | [...] |
| nonroot          | false | []   | MustRunAs | MustRunAsNonRoot | [...] |
| privileged       | true  | [*]  | RunAsAny  | RunAsAny         | [...] |
| restricted       | false | []   | MustRunAs | MustRunAsRange   | [...] |

# Security Context Constraints – Summary

- Integration of SELinux is great benefit
- Do not use existing Security Context Constraints except:
  - restricted
  - nonroot
- Create dedicated SCC's with least privilege principle if necessary

Demo 6: <https://asciinema.org/a/280392>

Source: [13]



# Conclusion & Future Work

## Conclusion & Future Work

- OpenShift raises the bar by it's defaults, but must be further adjusted
- Quick releases with feature extension/adjustment challenges the security research
- Multiple components are dedicated developed by RedHat and are not spread for the community
- OpenShift 4.2 is already available and components have been refactored and, new features and concepts are available



**Thanks for you Attention  
and take care of your  
defaults ;)**

# References

- [1] <https://blog.openshift.com/introducing-red-hat-openshift-4-2-developers-get-an-expanded-and-improved-toolbox/>
- [2] <https://twitter.com/bradgeesaman/status/1188541320626921473>
- [3] <https://blog.openshift.com/red-hat-chose-kubernetes-openshift/>
- [4] <https://kubernetes.io/blog/2015/04/borg-predecessor-to-kubernetes/>
- [5] <https://ai.google/research/pubs/pub43438>
- [6] <https://twitter.com/iancoldwater/status/1174793280011325456>
- [7] <https://goo.gl/TNRxtd>
- [8] [https://troopers.de/downloads/troopers19/TROOPERS19\\_AR\\_VXLAN\\_Security.pdf](https://troopers.de/downloads/troopers19/TROOPERS19_AR_VXLAN_Security.pdf)
- [9] <https://kubernetes.io/docs/concepts/services-networking/network-policies/#default-deny-all-ingress-traffic>
- [10] [https://docs.openshift.com/container-platform/3.11/admin\\_guide/managing\\_networking.html#admin-guide-limit-pod-access-egress](https://docs.openshift.com/container-platform/3.11/admin_guide/managing_networking.html#admin-guide-limit-pod-access-egress)
- [11] [https://docs.openshift.com/container-platform/3.11/admin\\_guide/managing\\_networking.html#admin-guide-limit-pod-access-egress-router](https://docs.openshift.com/container-platform/3.11/admin_guide/managing_networking.html#admin-guide-limit-pod-access-egress-router)
- [12] <https://kubernetes.io/docs/concepts/services-networking/dns-pod-service/#pod-s-dns-policy>
- [13] <https://cookbook.openshift.org/users-and-role-based-access-control/how-can-i-enable-an-image-to-run-as-a-set-user-id.html>